

4 - Refresher on linear regression

In this practice lab, you will fit the linear regression parameters (w, b) to your dataset.

- The model function for linear regression, which is a function that maps from x (city population) to y (your restaurant's monthly profit for that city) is represented as

$$f_{w,b}(x) = wx + b$$

- To train a linear regression model, you want to find the best (w, b) parameters that fit your dataset.
 - To compare how one choice of (w, b) is better or worse than another choice, you can evaluate it with a cost function $J(w, b)$
 - J is a function of (w, b) . That is, the value of the cost $J(w, b)$ depends on the value of (w, b) .
 - The choice of (w, b) that fits your data the best is the one that has the smallest cost $J(w, b)$.
- To find the values (w, b) that gets the smallest possible cost $J(w, b)$, you can use a method called **gradient descent**.
 - With each step of gradient descent, your parameters (w, b) come closer to the optimal values that will achieve the lowest cost $J(w, b)$.
- The trained linear regression model can then take the input feature x (city population) and output a prediction $f_{w,b}(x)$ (predicted monthly profit for a restaurant in that city).

5 - Compute Cost

Gradient descent involves repeated steps to adjust the value of your parameter (w, b) to gradually get a smaller and smaller cost $J(w, b)$.

- At each step of gradient descent, it will be helpful for you to monitor your progress by computing the cost $J(w, b)$ as (w, b) gets updated.
- In this section, you will implement a function to calculate $J(w, b)$ so that you can check the progress of your gradient descent implementation.

Cost function

As you may recall from the lecture, for one variable, the cost function for linear regression $J(w, b)$ is defined as

$$J(w, b) = \frac{1}{2m} \sum_{i=0}^{m-1} (f_{w,b}(x^{(i)}) - y^{(i)})^2$$

- You can think of $f_{w,b}(x^{(i)})$ as the model's prediction of your restaurant's profit, as opposed to $y^{(i)}$, which is the actual profit that is recorded in the data.
- m is the number of training examples in the dataset

Model prediction

- For linear regression with one variable, the prediction of the model $f_{w,b}$ for an example $x^{(i)}$ is represented as:

$$f_{w,b}(x^{(i)}) = wx^{(i)} + b$$

This is the equation for a line, with an intercept b and a slope w

Implementation

Please complete the `compute_cost()` function below to compute the cost $J(w, b)$.

Exercise 1

Complete the `compute_cost` below to:

- Iterate over the training examples, and for each example, compute:
 - The prediction of the model for that example

$$f_{wb}(x^{(i)}) = wx^{(i)} + b$$

- The cost for that example

$$cost^{(i)} = (f_{wb} - y^{(i)})^2$$

- Return the total cost over all examples

$$J(w, b) = \frac{1}{2m} \sum_{i=0}^{m-1} cost^{(i)}$$

- Here, m is the number of training examples and \sum is the summation operator

If you get stuck, you can check out the hints presented after the cell below to help you with the implementation.

6 - Gradient descent

In this section, you will implement the gradient for parameters w , b for linear regression.

As described in the lecture videos, the gradient descent algorithm is:

$$\begin{aligned} &\text{repeat until convergence: } \{ \\ &\quad b := b - \alpha \frac{\partial J(w, b)}{\partial b} \\ &\quad w := w - \alpha \frac{\partial J(w, b)}{\partial w} \\ &\} \end{aligned} \tag{1}$$

where, parameters w , b are both updated simultaneously and where

$$\frac{\partial J(w, b)}{\partial b} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{w,b}(x^{(i)}) - y^{(i)}) \tag{2}$$

$$\frac{\partial J(w, b)}{\partial w} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{w,b}(x^{(i)}) - y^{(i)})x^{(i)} \tag{3}$$

- m is the number of training examples in the dataset
- $f_{w,b}(x^{(i)})$ is the model's prediction, while $y^{(i)}$, is the target value

You will implement a function called `compute_gradient` which calculates $\frac{\partial J(w)}{\partial w}$, $\frac{\partial J(w)}{\partial b}$

Exercise 2

Please complete the `compute_gradient` function to:

- Iterate over the training examples, and for each example, compute:
 - The prediction of the model for that example

$$f_{wb}(x^{(i)}) = wx^{(i)} + b$$

- The gradient for the parameters w , b from that example

$$\begin{aligned} \frac{\partial J(w, b)}{\partial b}^{(i)} &= (f_{w,b}(x^{(i)}) - y^{(i)}) \\ \frac{\partial J(w, b)}{\partial w}^{(i)} &= (f_{w,b}(x^{(i)}) - y^{(i)})x^{(i)} \end{aligned}$$

- Return the total gradient update from all the examples

$$\frac{\partial J(w, b)}{\partial b} = \frac{1}{m} \sum_{i=0}^{m-1} \frac{\partial J(w, b)}{\partial b}^{(i)}$$

$$\frac{\partial J(w, b)}{\partial w} = \frac{1}{m} \sum_{i=0}^{m-1} \frac{\partial J(w, b)}{\partial w}^{(i)}$$

- Here, m is the number of training examples and \sum is the summation operator

If you get stuck, you can check out the hints presented after the cell below to help you with the implementation.